

GaitNet: Greedy, Acyclic Quadruped Gait Generation

Owen Sullivan



A Dissertation
Submitted to the Faculty
of the
WORCESTER POLYTECHNIC INSTITUTE
in partial fulfillment of the requirements for the
Degree of Master of Science
in
Robotics Engineering
December 2025

APPROVED:

Professor Mahdi Agheli
Advisor
Worcester Polytechnic Institute

Professor Jing Xiao
Committee Member
Worcester Polytechnic Institute

Professor Constantinos Chamzas
Committee Member
Worcester Polytechnic Institute

Abstract

Quadruped locomotion across unstructured terrain remains a longstanding challenge in robotics, requiring controllers that can adapt dynamically to complex and unpredictable environments. Traditional approaches rely on model predictive control or pre-defined gait schedules, offering stability but limited adaptability. Conversely, end-to-end learning methods enable agile and versatile behaviors but often sacrifice interpretability, safety, and real-time feasibility. This thesis introduces *GaitNet*, a hybrid control framework that integrates a greedy, neural network-based gait planner with a traditional model-based controller to generate dynamic, acyclic locomotion for quadruped robots.

The proposed system consists of two primary components: a *Footstep Evaluation Network* that learns terrain-aware footstep cost maps, and GaitNet, a reinforcement learning-based gait selector that ranks and executes feasible footstep actions in real time. Together, these modules enable a balance between the adaptability of learning-based control and the robustness of analytical motion planning. Trained in NVIDIA Isaac Lab using parallel GPU simulation, the system is evaluated across a range of terrain difficulties and commanded velocities.

Experimental results show that GaitNet achieves a 77.7% mean survival rate in challenging terrain, outperforming a single-leg motion planner baseline by more than a factor of two. Ablation studies further reveal that dynamic swing duration offers limited benefit in static environments, while removing pre-computed footstep costs leads to improved performance and more efficient gait patterns. These findings highlight the advantages of direct policy learning for coordinated, multi-leg motion without over-reliance on heuristic priors.

Overall, this work demonstrates that greedy, neural network-based planning can produce dynamic, non-gaited quadruped motion. The presented framework bridges the gap between reactive learning and structured control, providing a foundation for future research toward fully autonomous, terrain-adaptive legged locomotion in real-world settings.

Contents

List of Figures

List of Tables

List of Abbreviations

CNN	Convolutional Neural Network
DQN	Deep Q Network
MCTS	Monte Carlo Tree Search
ML	Machine Learning
MLP	Multi-Layer Perceptron
MPC	Model Predictive Control
PPO	Proximal Policy Optimization
RL	Reinforcement Learning

List of Symbols

Symbol ¹	Description	Units ²
$\mathbf{f}_c$	Footstep candidate Vector of <i>leg index</i> (or -1 for no leg), dx , dy , and <i>cost</i> . With dx and dy being offsets from the nominal foot position in the base frame.	$[- \text{ m } \text{ m } -]^T$
$\mathbf{f}_a$	Footstep action Vector of <i>leg index</i> (or -1 for no leg), dx , dy , and <i>swing duration</i> . With dx and dy being offsets from the nominal foot position in the base frame.	$[- \text{ m } \text{ m } \text{ s}]^T$
$\mathbf{r}_w$	COM position in world frame	m
$\mathbf{q}_{rw}$	Root orientation (quaternion) in world frame	-
$\mathbf{v}_b$	Root linear velocity in base frame	m s^{-1}
ω_b	Root angular velocity in base frame	rad s^{-1}
$\mathbf{p}_{ib}$	End-effector i position in base frame	m
$\mathbf{g}$	Gravity vector in base frame	m s^{-2}
$\mathbf{c}_i$	End-effector i contact state	-
$\mathbf{q}$	Joint positions	-
$\dot{\mathbf{q}}$	Joint velocities	-
$\ddot{\mathbf{q}}$	Joint accelerations	-
τ	Joint torques	N m
$\mathbf{u}$	Control Input Vector of v_x , v_y , and ω_z	$[\text{m s}^{-1} \text{ m s}^{-1} \text{ rad s}^{-1}]^T$

¹A symbol defined with a subscript i , but displayed without it indicates the concatenation of all i elements, e.g. $\mathbf{c} = [\mathbf{c}_1 \ \mathbf{c}_2 \ \mathbf{c}_3 \ \mathbf{c}_4]^T$

²Units marked with "-" indicate either no unit or multiple units

Terminology

Within legged robotics, there are many conflicting or ambiguous terms. For clarity, the exact meanings of specific terms used in this paper are described below.

Term	Definition
Footstep planning	The process of selecting when and where to step.
Greedy	In the context of footstep planning, a planner which only provides instantaneous actions.
Dynamic locomotion/gait	The opposite of quasi-static locomotion; where balance is maintained through active control and dynamic interactions, not slow and stable placements. Sometimes as an adjective, dynamic (as in the title).
Acyclic gait	Sometimes "free-gait" or "non-gaited". Refers to a footstep plan which cannot be represented with phase cycles or transitions between gait primitives. The opposite of gaited or gait-based locomotion.

1 Introduction

Quadruped locomotion in unstructured environments remains a significant challenge, particularly when traditional gait cycles prove inadequate [?]. While many existing systems rely on periodic gait patterns or centralized planners [?], recent advances in learning-based methods have enabled more adaptive, non-gaited approaches [?]. This study proposes a novel method for generating non-gaited, dynamic footstep plans using a neural network-based greedy planner, aiming to achieve a balance between computational efficiency and dynamic performance.

1.1 Motivation and Vision

Legged robots hold significant promise for enabling mobility in environments that are inaccessible to wheeled and tracked systems. Their ability to traverse uneven, unstructured, and unpredictable terrains—such as disaster zones, staircases, dense vegetation, and rocky landscapes—positions them as key enablers for future exploration, inspection, and rescue missions.

Realizing this potential requires a robot to continuously and intelligently decide where and when to place its feet—a process known as *contact planning*. Unlike periodic gait patterns suitable for flat or predictable surfaces, real-world locomotion often demands dynamic, acyclic, and adaptive footstep sequences. Even minor errors in contact selection can lead to instability or failure, whereas well-chosen footholds enable efficient, agile, and robust movement.

The generation of such contact plans in real time remains a formidable challenge. It requires the system to perceive its surroundings, anticipate the outcomes of possible actions, and select feasible contact points within tight temporal constraints. Existing methods often fall into one of three categories, fully using traditional control—model predictive controllers, PD controllers, optimization—using neural networks, or a mix of the two: hybrid methods. These three methods offer different trade offs between deliberative, computationally intensive planning and fast, reactive control that lacks stability or generalization.

The vision of this research is to develop a hybrid control architecture that optimizes these trade offs. Specifically, this work aims to design a system capable of producing non-gaited, dynamic footstep plans that retain both the responsiveness of learning-based approaches and the reliability of model-based control. The objective is to provide quadruped robots with the decision-making agility necessary for navigating complex terrains while maintaining the stability and robustness demanded by real-world operation.

Through this approach, the thesis seeks to contribute toward a new class of hybrid locomotion frameworks that enable real-time, adaptive, and physically consistent quadruped mo-

tion—closing the gap between high-level perception and low-level control in legged robotics.

1.2 Research Gap

Quadruped control pipelines can be categorized into several approaches, many of which can be broadly classified as traditional, neural network-based, or hybrid methods.

Traditional approaches typically rely on model predictive control (MPC) frameworks for motion optimization, supported by lower-level controllers to track desired trajectories [?]. Many aspects of quadruped locomotion—such as joint torques and ground reaction forces—are smooth and well-suited to continuous optimization techniques [?]. However, these methods become less tractable when discrete contact states are introduced into the optimization problem [?]. To manage this complexity, most implementations either assume pre-defined gait sequences [?, ?] or encode discrete contact decisions within larger optimization structures [?]. While effective in structured environments, these strategies constrain the robot’s ability to exploit the full versatility of legged locomotion by reducing the dimensionality of expressible states for the system.

Neural network-based methods have emerged as an alternative, particularly with advances in deep reinforcement learning. These approaches often replace the MPC component of traditional control pipelines, learning to directly map observed states to joint torques or motion commands [?]. In doing so, they inherently address the mixed-integer nature of contact planning, allowing the network to learn both continuous and discrete aspects of locomotion simultaneously [?]. Despite their flexibility, such methods generally lack formal guarantees of stability, require substantial training data, and often struggle to generalize beyond their training distribution [?].

Hybrid approaches, which integrate learning-based modules within model-based control frameworks, have recently gained attention as a potential middle ground. In these systems, neural networks are typically employed to solve specific subproblems—such as dynamically switching between fixed gaits or greedily selecting individual footholds—while the overall motion execution remains governed by a traditional controller [?, ?]. This division of responsibilities enables learning to focus on the non-smooth or combinatorial components of locomotion, while preserving the interpretability, safety, and robustness inherent to model-based systems.

Looking at how each of these categories handles the contact planning problem, hybrid control schemes have demonstrated promise, the challenge of contact and footstep planning continues to limit their performance. Traditional optimization-based planners struggle with the discrete nature of contact transitions, resulting in high computational costs that preclude real-time operation [?]. Alternatively, many rely on pre-specified gait patterns [?, ?, ?, ?], constraining adaptability to irregular or unpredictable terrain. Even more advanced approaches, such as Monte Carlo Tree Search (MCTS)-based planners, have achieved impressive dynamic behaviors but remain computationally demanding [?, ?]. Finally, end to end learning methods solve these problems implicitly, but at the cost of control theory guarantees and interpretability.

Despite recent progress, a clear research gap remains. Current hybrid methods do not yet achieve the balance between adaptability and stability that would enable fully dynamic,

acyclic locomotion in complex environments. In particular, there is a lack of systems that can match the contact planning agility of neural network-based approaches while retaining the formal guarantees and reliability of traditional control frameworks [?, ?]. Recent supervised learning estimators have demonstrated the potential of machine learning for footstep evaluation; however, their design—often restricted to single-leg motions with fixed timing—still falls short of the dynamic, non-gaited behaviors exhibited by end-to-end learning systems.

This motivates a central research question:

Can a hybrid control pipeline, which integrates a greedy, neural-network-based planner with a traditional model-based controller, generate dynamic, acyclic gaits for robust quadruped locomotion in challenging environments?

1.3 Hypothesis

To address the identified research gap, we hypothesize that a hybrid control pipeline, incorporating a greedy, neural network-based footstep planner and a traditional model-based control framework, can effectively generate dynamic, non-gaited locomotion for quadruped robots in challenging environments. By extending continuous optimization capabilities to support multi-leg motion with dynamic swing durations, we believe that such a planner can produce footstep sequences that enable the robot to traverse complex terrains while maintaining both stability and efficiency.

1.4 Approach

The proposed approach bypasses traditional discrete sampling or candidate generation networks entirely. Instead, a novel single-network architecture, *GaitNet*, is used to evaluate and select the most appropriate action at each time step. By directly exploring the output space through projected gradient ascent, GaitNet optimizes spatial footstep coordinates directly over its own reward surface. By integrating this continuous optimization planner within a traditional model-based control framework, we harness the advantages of rapid learning-based evaluation and robust model-based stabilization, yielding an adaptable quadruped locomotion system.

1.5 Contributions

Quadruped locomotion in unstructured environments demands footstep planners that combine adaptability, computational efficiency, and physical consistency. This paper addresses the challenge of generating dynamic, acyclic contact sequences by exploring a hybrid control framework that integrates a greedy, neural network-based footstep planner—GaitNet—within a traditional model-based controller. Building on insights from recent learning-driven approaches while maintaining the reliability of structured control, this work aims to close the gap between flexible footstep selection and stable trajectory execution. The following contributions summarize the central advances introduced in this paper:

- Development of a novel greedy footstep planner utilizing a single neural network to evaluate and select continuous footsteps via projected gradient ascent, eliminating discrete sampling stages.
- Demonstration that the proposed planner can generate dynamic, acyclic gaits in challenging environments.
- Empirical evaluation on physical robots, showcasing Real-World locomotion capabilities on challenging terrain utilizing a Unitree Go1 quadruped mapped against contemporary state-of-the-art planners.

1.6 Thesis Structure

The remaining chapters are organized as follows.

Chapter 2 presents the background necessary to contextualize this work, covering fundamental concepts in quadruped locomotion, gait generation, and learning-based control methods. This chapter establishes the theoretical foundations on which the proposed approach is built.

Chapter 3 describes the methodology, detailing the design and implementation of the CNN-based greedy planner, GaitNet, including the underlying network architecture and training pipeline. It explains how this module integrates with a model-based control framework to enable dynamic footstep planning.

Chapter 4 reports the experimental results and evaluates the performance of the proposed planner across a range of challenging scenarios. This chapter also discusses the strengths and limitations of the method, examining how the system behaves under diverse conditions.

Chapter 5 concludes the thesis by summarizing the key findings and reflecting on the broader implications for quadruped locomotion research. It also identifies potential directions for future work, particularly avenues for improving adaptability, stability, and real-world deployment.

2 Background

Recent advances in legged locomotion have largely focused on either optimization-based motion planning or learning-based policy generation. The former, categorized as traditional control approaches, offers formal guarantees but faces significant computational hurdles when dealing with dynamic, non-gaited footstep selection. The latter, end-to-end neural networks, provides adaptability and implicitly solves the discrete contact problem, but lacks control theory guarantees and interpretability.

The most promising efforts to bridge this gap are hybrid methods, which integrate learning modules within model-based frameworks. Our proposed approach seeks to enhance this hybrid strategy by developing a computationally efficient, non-gaited footstep planner. This chapter reviews the foundations upon which our method builds: the limitations of traditional control and end-to-end Neural Networks, and the state-of-the-art in hybrid methods, including learned gait policies and explicit footstep planners. This review establishes the necessity of our novel hybrid architecture.

2.1 Traditional Control Approaches

For the purpose of this thesis, we define traditional control approaches as those that rely on explicit optimization of footstep positions, contact sequences, and timing, often through mixed-integer programming or search-based methods. These approaches provide guarantees associated with control theory but often suffer from high computational costs and limited real-time applicability, especially for methods which attack the non-gaited footstep planning problem.

Deits and Tedrake [?] formulate footstep placement as a mixed-integer quadratic program over convex terrain regions, enabling robust bipedal foothold selection. However, their method does not optimize contact timing and is limited to sequential foot placement, restricting agility and generalization to more complex contact sequences.

Extensions to this framework incorporate contact scheduling into the optimization itself. Winkler et al. [?] fully optimize contact duration and ordering, allowing multiple steps to be planned simultaneously rather than sequentially. While this increases flexibility, the resulting computation is too slow for real-time control. Similarly, Aceituno-Cabezas et al. [?] integrate footstep positions and gait timings into a mixed-integer convex program based on centroidal dynamics, achieving simultaneous contact and motion planning. Taouil et al. [?] further explore non-gaited footstep planning via MCTS, generating dynamic multi-step sequences in real time, though computation remains high due to the branching factor. Akizhanov et al. [?] improve MCTS efficiency using a learned classifier to prune contact configurations and

a target adjustment network to compensate for low-level control errors, but the approach is still computationally intensive.

These methods provide precise, globally consistent locomotion plans, but high computational costs and limited real-time applicability restrict their use in dynamic, acyclic gait scenarios.

2.2 End-to-End Neural Networks

End-to-end neural network approaches for legged locomotion directly map sensory inputs to control outputs, bypassing explicit modeling of dynamics or footstep planning.

Zhang et al. [?] propose an end-to-end RL approach mapping proprioceptive input directly to joint targets, reaching 2.5 m/s and generalizing across terrains via a terrain curriculum and curiosity rewards. However, such models require slow training, and re-training for new environments. Zhang et al. [?] address multi-gait transitions using heuristically defined tables between similar gaits, improving flexibility but still constraining behaviors to predefined families.

While these models achieve impressive feats, they often require extensive training and struggle to generalize to unseen terrains or tasks. This motivates approaches that combine the adaptability of learning-based methods with the structure and guarantees of traditional control, enabling fast, terrain-adaptive footstep planning without the need for extensive re-training.

2.3 Hybrid Methods

Hybrid methods for legged locomotion combine elements of traditional control (e.g., optimization or Model Predictive Control, MPC) and learning-based approaches to achieve fast, terrain-adaptive footstep generation. These methods often solve high-level planning problems with learned models, while delegating low-level control to traditional, high-frequency solvers. This division of labor enables real-time performance without sacrificing the expressiveness offered by deep learning.

2.3.1 Learned Policies for Gait and Trajectory Modulation

A major class of hybrid methods utilizes deep learning to generate agile, robust locomotion policies, often focusing on body motion and gait modulation rather than explicit, step-by-step foot placement. These approaches typically learn a mapping from sensor data to control commands, relying on fixed or implicit gait patterns.

Deep learning has shown strong potential for generating such robust policies. Shi et al. [?] modulate trajectory generator parameters in real time for energy-efficient walking, while Xie et al. [?] train Reinforcement Learning (RL) policies on centroidal dynamics models to output body accelerations under fixed gait heuristics. Other methods focus on blind, proprioceptive-based control: Duan et al. [?] learn step-to-step transitions using proprioception, and Siekmann et al. [?] achieve blind stair climbing through LSTM-based propriocep-

tive policies. Lee et al. [?] even infer terrain structure from proprioceptive history using a temporal CNN and automated curriculum learning.

Subsequent works extend these ideas through explicit high-level gait selection. Da et al. [?] use a Deep Q-Network (DQN) to choose among predefined gait primitives executed by a low-level controller, while Yang et al. [?] learn policies that output gait parameters—frequency, swing ratio, and phase offsets—interpreted by a phase integrator. Both enable efficient gait modulation but generally assume flat terrain and heuristic foot placement. Sun et al. [?] integrate offline gait optimization with contact-implicit trajectory optimization and high-frequency MPC, achieving dynamic control but limiting adaptability to a set of discrete, speed-dependent gaits.

Although effective at generating robust body motions, these methods primarily rely on fixed or implicit gait patterns and largely lack explicit, flexible control over individual foot-steps based on local terrain features.

2.3.2 Learning-Based Footstep Planners

In contrast to gait modulation, learning-based footstep planners focus on explicitly learning the footstep planning component. These methods combine perception and control through data-driven models, often coupling the learned components with downstream MPC or search-based planning to enforce dynamic feasibility.

Early methods demonstrated the concept but faced scope limitations. FootstepNet [?] generates heuristic-free plans via deep RL but is limited to bipeds in simple 2D environments. DeepLoco [?] introduces a hierarchical policy for footstep and motion planning, yet its coarse terrain input and biped focus limit dynamic adaptability. Other approaches separate the planning concerns: DeepGait [?] separates footstep planning from motion optimization, enabling modularity but relying on static gaits. RLOC [?] uses a convolutional encoder to pass terrain data embeddings to multiple RL modules before generating joint torques with a traditional whole body controller. One of these RL modules is a footstep planner, but the system uses a fixed gait schedule and the footstep planner is only evaluated once at the beginning of each cycle.

A key challenge is moving beyond fixed gaits while maintaining stability. A recent convolutional evaluator [?] uses a CNN to produce non-gaited footstep sequences with greedy selection, though it is constrained to one-leg-at-a-time motion and coarse foothold discretization. Omar et al. [?] and Villarreal et al. [?] use CNNs for terrain classification to inform MPC or heuristic-free selection, but both depend on fixed gait sequences.

Recent Advances in Terrain Generalization

More recent advances expand terrain generalization and multi-contact reasoning. Tolomei et al. [?] created a framework generalized to 3D terrain and variable foot shapes, ranking footholds based on terrain curvature, but still follow a fixed gait. Chen et al. [?] apply a DQN to a hexapod, simultaneously selecting both next foot positions and gait, with fallback to predefined behaviors outside expected states, but motion is constrained to three-leg-at-a-time gaits. Similarly, Yao et al. [?] output both swing leg motions and desired robot pose

from a deep RL policy but remain limited to a trotting gait. Meduri et al. [?] use a DQN to learn 3D foothold cost maps, but retain fixed biped gait timing.

Overall, these approaches highlight trade-offs between expressive, non-gaited footstep planning and computational efficiency. CNN-based and RL methods improve terrain generalization, but fixed gaits or sequential leg motions remain common constraints for achieving real-time performance and maintaining stability.

2.3.3 Greedy Planning and Learned Foothold Classification

To complement learning-based approaches, other hybrid methods focus on improving the efficiency of the planning step itself through greedy search or by leveraging learned terrain models for safe foothold identification.

Greedy Planners for Efficiency

Greedy planners offer computationally efficient alternatives to exhaustive search, often prioritizing goal-directed heuristics or local stability metrics. Gao et al. [?] propose GH-QP, a greedy-heuristic hybrid that incorporates expected robot speed into a footstep planning heuristic, though it is limited to bipedal systems and lacks support for dynamic footstep plans. Zucker et al. [?] use A* with a learned terrain cost and Dubins car heuristic for footstep planning, enabling effective search in SE(2) but restricted to static, gated locomotion. Kalakrishnan et al. [?] combine greedy terrain-cost search with a coarse-to-fine pose optimization pipeline, but rely on fixed gait patterns.

It is worth noting that greedy elements are already incorporated into several learning-based methods discussed previously, such as recent ML planners [?], DeepLoco [?], and Villarreal et al. [?]. While these works demonstrate the potential of greedy methods for efficiency, none explicitly support both dynamically unstable and non-gaited footstep generation.

Learned Foothold Classification

A related line of research focuses solely on determining safe stepping regions (foothold classification) using learned terrain models to support downstream planners that enforce safety constraints. Asselmeier et al. [?] generate steppability maps from visual inputs to guide a trajectory-optimization-based planner. Omar et al. [?] propose a two-stage classifier that first identifies steppable regions before selecting footholds, emphasizing safety. Similarly, Omar et al. [?] use a CNN to identify safe footholds, which are then passed to an MPC for real-time planning, but the approach relies on fixed gait sequences.

These methods provide robust terrain understanding but are typically designed as standalone perception modules, separate from the full dynamic footstep generation process, and thus are not directly applicable to dynamic or non-gaited motion planning without significant integration.

2.4 Synthesis

The preceding review of quadruped locomotion control, encompassing traditional optimization-based methods, end-to-end neural networks, and hybrid control schemes, reveals a fundamental trade-off: control theory guarantees and robustness versus computational efficiency and adaptability to unstructured, non-gaited motion. Traditional methods provide stability but suffer from high computational cost, particularly when attempting to solve the mixed-integer programming problem inherent in non-gaited contact planning [?, ?]. Conversely, end-to-end learning approaches achieve impressive dynamic, acyclic behaviors but lack formal guarantees and generalizability beyond their training distribution [?].

Hybrid control architectures have emerged as the promising middle ground, integrating learned modules within model-based frameworks to balance agility and stability. However, existing hybrid footstep planners, such as recent ML planners [?], are often constrained by enforcing sequential, one-leg-at-a-time motion or relying on fixed gait timing. These constraints fundamentally limit the robot’s ability to execute the fully dynamic, acyclic locomotion required for optimal performance in complex terrain.

This thesis addresses the identified gap by proposing a novel hybrid architecture built upon the efficiency of greedy, NN-based selection while overcoming the single-leg and fixed-timing limitations of prior work. The GaitNet planner is designed to simultaneously evaluate and select footstep candidates for multiple legs with dynamic swing durations. By delegating the computationally intensive, discrete decision-making of contact selection to this neural network and feeding the resulting plans into a robust, model-based controller, we aim to achieve a system that combines the agility and non-gaited expressiveness of learning-based approaches with the stability and physical consistency of traditional control frameworks. This synthesis validates the necessity of a new hybrid approach capable of supporting dynamic, multi-leg contact planning for robust quadruped locomotion.

3 Methodology

3.1 System Overview

This work presents a control framework for quadruped robots capable of generating dynamic, acyclic gaits in challenging environments. The framework employs a hierarchical architecture that processes robot state and terrain data to produce footstep commands for the robot controller.

The first stage of the framework involves the user input $\mathbf{v}^{\text{usr}}$ being supplied alongside the robot state to a gait generation network (*GaitNet*). *GaitNet* acts as an evaluative surface, performing continuous optimization over the physical workspace utilizing projected gradient ascent. Unlike discrete sampling methods, *GaitNet* ranks these continuous outputs iteratively, producing the optimal footstep action $\mathbf{f}_a$, which is passed to the robot controller. The robot controller is the MIT Mini-Cheetah Controller [?], implemented in python by [?].

This hierarchical design enables strict control over the range of possible actions the robot can execute. Between the two stages, candidate actions are filtered to ensure that only valid, prescribed movements are permitted.




Figure 3.1: Block diagram of the proposed framework. Novel aspects shown in blue. The user defines an input direction $\mathbf{v}^{\text{usr}}$, which *GaitNet* (??) uses along with the robot state $\mathbf{x}_g$ and the current foot positions $\mathbf{p}_f$ to evaluate candidate footsteps. Instead of discrete sampling, *GaitNet* utilizes a continuous projected gradient ascent optimization strategy, generating the optimal footstep action $\mathbf{f}_a$ to send to the robot controller. The *MIT Mini-Cheetah Controller* [?] receives $\mathbf{f}_a$ and executes the corresponding actions.

3.2 Simulation Environment

The simulation environment used for this project is NVIDIA Isaac Lab [?]. This framework was selected for several reasons. It is a modern platform with a Python interface designed specifically for reinforcement learning applications and GPU parallelism. The use of GPU parallelism enables significantly faster simulation and data collection, albeit at the cost of increased programming complexity and reduced compatibility with older hardware. Furthermore, the extensive collection of example and community projects provides valuable references for implementing the simulation features required in this work.

Although both simulation and learning processes are executed on the GPU, the robot controllers operate on the CPU. ?? illustrates the overall processing flow. The simulation environment runs entirely on the GPU, where multiple robots are simulated in parallel. Meanwhile, the robot model predictive controllers (MPCs) execute on the CPU, with each controller running concurrently. The CPU and GPU communicate at every simulation step to exchange robot states and corresponding control actions.




Figure 3.2: Block diagram showing the programming tasks computed on the CPU versus the GPU. The full simulation is run in parallel using NVIDIA Isaac Lab on the GPU, while the robot MPCs are run in parallel on the CPU. Multiple Cores in the CPU and multiple "MPC & PD" in each core show the hierarchy of how the parallel computing workload is assigned.

NVIDIA Isaac Lab uses a declarative system of Python dataclasses to define the simulation environment. For this work, a custom environment is used (??), including:

- Unitree Go 1—Configured to use force control for each joint.
- Terrain raycasts—Measure the height of the terrain at each possible footstep location. This emulates the lidar processing steps of a full vision pipeline.

⁰As is explained in later sections, GaitNet only outputs one footstep action at a time. The two separate foot target locations shown in the GaitNet visual would be selected sequentially with a 1/25s delay. This visual is meant to highlight the near simultaneous nature of the footstep selection process.

- Custom terrain—A grid of 12×12 sub-terrains (4×4 m each) with increasing difficulty (??). Each sub-terrain consists of a 12 cm square grid pattern with missing sections. The void density varies from approximately 0% to 40%. The quantity of terrain is limited by GPU memory.




Figure 3.3: Image of a Unitree Go 1 navigating the terrain. Red region shows area converted into heightmap for front left leg. Black regions show voids in the terrain. Green arrow shows desired velocity vector.




Figure 3.4: Image showing the full terrain used in simulation. Full terrain is a grid of 12×12 sub-terrains with increasing difficulty. Sub-terrains are 4×4 m grids with missing sections.

3.3 GaitNet

GaitNet is designed to isolate the optimal footstep action from a continuous configuration space using gradient-based optimization. Instead of selecting from a discrete grid populated by a prior evaluation network, GaitNet serves as the primary reward estimator, enabling direct optimization in the continuous output space.

Specifically, GaitNet optimizes the one-hot encoded leg index $\mathbf{f}_c'$ and the spatial coordinates via projected gradient ascent. The inclusion of a no-action candidate, evaluated directly alongside the optimized paths, allows the greedy generation of acyclic gaits in which multiple feet can be in the swing phase simultaneously.

3.3.1 Architecture

GaitNet is responsible for ranking footstep candidates based on the one-hot encoded leg index $\mathbf{f}_c'$ and the GaitNet input $\mathbf{x}_g$:

$$\mathbf{x}_g = \begin{bmatrix} \mathbf{p}_{b,xy} \\ \mathbf{r}_{w,z} \\ \mathbf{v}_b \\ \omega_b \\ \mathbf{u} \\ \mathbf{g} \\ \mathbf{c} \end{bmatrix} \quad (3.1)$$

Descriptions of the first five terms dictate the primary robot state, tracking current base position $p_{b,xy}$ and rotation $r_{w,z}$ in world coordinates, body linear $\mathbf{v}_b$ and angular ω_b velocities, and current foot positions $\mathbf{u}$. The additional terms are $\mathbf{g}$ and $\mathbf{c}$, being the gravity vector and contact state respectively. Using $\mathbf{x}_g$ and $\mathbf{f}_c'$, GaitNet outputs a logit associated with each evaluated spatial coordinate and the desired swing duration if that step is selected.




Figure 3.5: GaitNet architecture. Sections of the diagram include the footstep candidate encoder in red, robot state encoder in blue, and shared trunk in green. The final output is a logit encoding the value of this option and the desired swing duration if this action is taken.

The model is designed to jointly evaluate robot state and footstep candidates with a middle fusion architecture [?]. The architecture (shown in ??) consists of two encoders, a shared trunk, and two task-specific output heads:

- Robot state encoder: The robot state vector is processed by a three-layer feedforward

network with intermediate dimensionality of 128 and ReLU activations. This produces a fixed-dimensional latent representation of the current robot state.

- Footstep encoder: Each footstep candidate is encoded by a two-layer feedforward network with hidden size 64 and ReLU activations. No-action candidates are represented through a fixed embedding.
- Shared trunk: The concatenated robot state and footstep embeddings are processed by a three-layer feedforward trunk, reducing dimensionality while applying layer normalization and ReLU nonlinearity.
- Output heads: Two parallel prediction heads are applied to the shared trunk:
 - Logit head: A single-layer network outputs the predicted reward value as a logit.
 - Duration head: A single-layer network with sigmoid activation outputs a normalized swing duration, which is then scaled to the range (0.1, 0.3) s.

This design allows the network to assign an expected value and a feasible swing duration to any given spatial coordinate, while explicitly supporting a no-action option via a dedicated embedding.

3.3.2 Projected Gradient Ascent Optimization

Instead of sampling from a discrete pool of generated candidates, we formulate the footstep selection as a continuous optimization problem over the spatial coordinates of the footstep action. GaitNet’s reward prediction logit is treated as a differentiable objective function maximized via projected gradient ascent.

For each foot, $N = 4$ random initial spatial samples are drawn uniformly from the reachable kinodynamic workspace. For each sample, the spatial coordinates are updated in parallel for 4 steps to maximize the expected value output logit computed by GaitNet. After each gradient update step, the new spatial coordinates are projected back to the nearest valid configuration within the reachable stepping continuous space (cspace) to ensure physical validity. This forces the generated footsteps to avoid un-steppable terrain like holes while preserving maximum reward.

Let the fixed ”no-action” embedding be evaluated alongside the spatial updates. After the 4 projection and ascent steps, the candidate with the highest reward logit across all 4×4 options (4 tracks, 4 legs) is compared to the predicted value of the ”no-action” candidate. The best resulting footstep or no-action command is then deployed to the inner controller. This optimization runs at control frequencies, tightly integrating flexible learning-driven evaluation with robust multi-leg acyclic gait coordination.

3.3.3 Training

Due to the inclusion of the $\mathbf{c}$ term in the input, GaitNet cannot be effectively trained using supervised learning because of the high dimensionality of the problem. Instead, it is trained using PPO [?]. In this setup, the actor network is GaitNet (??), while the critic follows a

similar architecture but omits the duration head. All logits from the critic are passed through a final MLP to produce a single value estimate. This MLP consists of two hidden layers of size 64, with ReLU activations applied to all layers except the output.

A custom actor-critic implementation was necessary to handle GaitNet’s dual outputs. The logits define a categorical distribution over footstep candidates, while the duration output is treated as a separate normal distribution with a fixed standard deviation of 0.01 s for each action. To prevent multiple no-action candidates from skewing the selection, all but one no-action candidate are removed (set to $-\infty$).

The total log probability of an action is computed as the sum of the categorical log probability of the selected footstep candidate and the log probability of the duration under the normal distribution. During action sampling, a footstep candidate is first sampled from the categorical distribution, followed by sampling the duration from its associated normal distribution. The footstep candidate and duration are then combined to form a complete footstep action.

Following the reward structure, termination conditions, commands, and hyperparameters outlined in ??, GaitNet is trained for 700 environment steps using 100 agents in parallel. Each environment step simulates 20 s. The 700 step end point was chosen because improvement drastically slowed by then.

4 Results and Discussion

4.1 GaitNet

The primary objective of GaitNet is to generate fully dynamic and acyclic gaits for quadruped robots. An example swing sequence is shown in ??, illustrating a acyclic, dynamic gait in which the robot performs motions with up to two feet off the ground simultaneously. Swing durations are non-uniform, with some legs remaining in the swing phase longer than others. A video of GaitNet navigating challenging terrain can be seen [here](#).

Examining ??, a short step is observed for the front-left leg (green) at 3s, while the rear-right leg (red) executes a longer step at 5s. Additionally, between 0–7.5s, the robot is commanded to follow a slow input, moving one foot at a time. After 7.5s, the input speed increases, prompting a more dynamic gait. During this faster phase, the robot occasionally executes two steps simultaneously and does not follow a strict alternating pattern.




Figure 4.1: Example swing schedule for a single gait cycle. Each row represents a leg, with color indicating the leg is in swing phase.

To ensure the consistency of the training process, GaitNet was trained three times. Due to computational limitations, the 700-step training curriculum needed to be split into multiple runs. The runs were typically 350 steps long. This has the side effect of resetting the distribution of terrain levels, which can cause discontinuities in certain training metrics. Resetting the simulations did not appear to have an adverse affect on the learning process.

The mean episode reward during training is shown in ??. Each training instance shows a similar trend, with the reward increasing as the model learns to navigate the terrain.




Figure 4.2: Mean episode reward during GaitNet training. Each color represents a different training instance.

Another important metric is the mean steps per second, shown in ???. Initially, with limited training, the model moves its legs very frequently, leading to unstable motion. In initial experiments, it was found that this issue does not resolve well without guidance from the reward function. ??? provides more detail on this issue. The mean steps per second starts high and decreases as the model learns to take more stable steps.




Figure 4.3: Mean steps per second during GaitNet training. Each color represents a different training instance.

During training, the GaitNet model learns to adjust swing durations based on its input, as illustrated in ?? and ??. The mean swing duration settles to 0.24s with a standard deviation of 0.034s. ?? shows exactly what has been learned more clearly. For this histogram all

swing durations from a GaitNet test described in ?? were recorded. We can see a large peak between 0.225s and 0.25s, with a significant amount of swings down to 0.125s.




Figure 4.4: Mean swing duration across all environments during GaitNet training. Each color represents a different training instance.




Figure 4.5: Swing duration standard deviation across all environments during GaitNet training. Each color represents a different training instance.




Figure 4.6: Histogram of swing durations selected by GaitNet.

4.1.1 Discussion

The results presented demonstrate that GaitNet successfully generates dynamic, acyclic gaits for quadruped locomotion in challenging environments. The system can synthesize a wide variety of actions, and the results from the swing schedule analysis indicate that the system is capable of dynamically varying step timing. The system’s response to changing control inputs—taking a cautious approach at low speeds and executing simultaneous steps at high speeds—highlights the system’s flexibility in generating acyclic gaits without pre-defined patterns.

Regarding the training metrics, there was no specific target for an ideal steps per second, but the change during training shows that the model is correctly learning when and when not to step. Similarly, there was no specific target value for the swing duration metrics. They are presented to verify that GaitNet is capable of learning to utilize its swing duration output. The convergence of the mean swing duration and standard deviation indicates that the model uses a wide range of its allowable swing durations. The histogram further clarifies this, showing that while the system favors durations between 0.225 s and 0.25 s for normal motion, it utilizes durations down to 0.125 s for situational motions, such as moving across tricky terrain or recovering from falls. These figures confirm that the model learns to effectively utilize dynamic swing durations to adapt to varying conditions.

4.2 Physical Robot Experiments and Baseline Comparison

To evaluate GaitNet’s continuous optimization strategy in unstructured real-world scenarios, a baseline method is established utilizing the *Legged Perceptive Stack* [?]. This state-of-the-art framework employs Nonlinear Model Predictive Control (NMPC) and Whole-Body Control (WBC) based on OCS2 [?].

In order to benchmark GaitNet and its specific structural ablations against this baseline, a custom physical test environment was constructed mirroring the simulation parameters

(??). Hardware experiments are conducted utilizing a Unitree Go1 quadruped. The physical terrain consists of narrow strips of ground interspersed with gaps, characterized by varying terrain difficulty metrics defined as the percentage of unsteppable terrain (ranging from 0% to 40% difficulty).

We systematically evaluate the success rate, total steps taken, and overall traversal time of the following methods across 50 physical trials of 20s each:

1. **Legged Perceptive Stack:** The model predictive control baseline [?].
2. **GaitNet-Single:** A targeted ablation of GaitNet restricted to only one leg in swing at any given time.
3. **GaitNet-Trot:** A targeted ablation of GaitNet restricted to only opposing legs in swing concurrently.
4. **GaitNet (Full):** The proposed unlimited acyclic footprint generation model.

An attempt is considered successful if the robot completes the traversal without falling or stepping into the structurally removed terrain sections (invalid terrain).




Figure 4.7: Unitree Go1 navigating the physical test environment used for baseline comparison. The terrain consists of narrow strips with missing sections mapping to varying terrain difficulties (percentage of unsteppable terrain).

4.2.1 Physical Locomotion Results

?? and ?? summarize the performance of GaitNet (Full and ablations) alongside the baseline method across a range of terrain difficulties spanning from flat ground to 40% unsteppable gaps.

Table 4.1: Placeholder: Real-World Metrics on Unitree Go1

Method	Success Rate	Avg. Traversal Time [s]	Avg. Steps Taken
Legged Perceptive Stack	XX%	XX.X	XXX
GaitNet-Single	XX%	XX.X	XXX
GaitNet-Trot	XX%	XX.X	XXX
GaitNet (Full)	XX%	XX.X	XXX




Figure 4.8: Hardware Evaluation Success Rates over terrain difficulties. GaitNet (Full) achieves superior fault-free traversal spanning challenging obstacle densities compared to restrictive gaits and the perceptive baseline.

4.2.2 Discussion

GaitNet’s unconstrained performance relative to the Legged Perceptive Stack and fixed-gait ablations reveals the clear advantage of dynamic, acyclic spatial optimization. The system’s ability to coordinate multi-leg motions and adapt continuous spatial timing concurrently minimizes the total steps required while avoiding invalid terrain. This improvement is particularly pronounced at higher terrain difficulties (e.g., 30-40% gap densities), where restrictive trot or single-leg strategies inherently clash with the unstructured valid footprint areas.

4.3 Swing Duration Ablation Study

In this section we present an ablation study to assess the impact of dynamic swing duration on GaitNet’s performance. We compare two models: the standard GaitNet formulation as described in ??, and a Duration-Ablated-GaitNet trained with a fixed swing duration. In the GaitNet formulation described previously, the model explicitly outputs a swing duration to use for every action. The corresponding swing duration is then used for the chosen action. For this ablation, we ignore the swing duration output of GaitNet and directly use a fixed value, 0.24s. This value was chosen because it is roughly the mean step duration used by

GaitNet. This will provide information about how important the variable swing duration is for this use case.

Now, we train the Duration-Ablated-GaitNet model, setting the swing duration to a constant 0.24s. The training process is identical to that of the standard GaitNet model, ensuring a fair comparison. The performance of both models is then evaluated using the same metrics outlined in ??.




Figure 4.9: Evaluation of Duration-Ablated-GaitNet across various terrain difficulties and commanded velocities. Overall survival rate of 67.2%. Mean success rate measured as the percentage of 50 episodes which completed 20s without terminating, under the termination conditions described in ??.

?? presents the performance of the Duration-Ablated-GaitNet. The results indicates negligible performance changes over standard GaitNet formulation. GaitNet was able to achieve an overall survival rate of 69.4% (??), while the Cost-Ablated-GaitNet achieved a survival rate of 67.2% (??). The mean episode reward during training, as shown in ??, shows a similar trend.




Figure 4.10: Mean episode reward during training for GaitNet and Duration-Ablated-GaitNet.

4.3.1 Discussion

The swing duration ablation study yields a surprising result: removing dynamic swing duration control has minimal impact on performance. The steady state behavior of the reward curves for both models indicates that the Duration-Ablated-GaitNet is able to learn a policy that performs comparably to the standard GaitNet, despite the lack of dynamic swing duration. This finding suggests two possibilities. First, the current training environment may not sufficiently challenge the system to exploit variable swing timing—the static terrain and relatively low speed requirements may not create scenarios where dynamic timing provides substantial benefits. Second, the reward function may not adequately incentivize optimal swing duration selection, causing the network to default to near-constant timing regardless of capability.

The observation that trained GaitNet naturally converges to a mean swing duration of 0.24 s with only modest variation (standard deviation of 0.034 s) supports this interpretation. These findings suggest that dynamic swing duration may not be a critical factor for GaitNet’s performance in the tested environments. However, it is important to note that this conclusion may not hold in more complex or dynamic environments, where the ability to adjust swing duration could provide a significant advantage. Future work in more dynamic environments or with revised reward structures may reveal greater utility for this feature.

4.4 Action Cost Ablation Study

In this section we present an ablation study to assess the impact of the footstep candidate cost on GaitNet’s performance. We compare two models: the standard GaitNet formulation as described in ??, and a Cost-Ablated-GaitNet trained with the the footstep candidate ($\mathbf{f}_c$) zeroed out in the input. In the GaitNet formulation described previously, we pass the estimated heuristic cost of a footstep action into the model ($\mathbf{f}'_c$ in ??), that is, the expected cost of taking the action if it were the only motion the robot was performing. For this

ablation study we replace that value with a zero during the training and evaluation process to see how it affects the network without changing its total parameters.

During training, the GaitNet model quickly learns to correlate the lower cost candidates with higher value logits. After this initial learning phase, the model learns to refine its predictions based on other features in the input. This is illustrated in ??.

The inspiration for this ablation study comes from observing that the footstep candidate cost and GaitNet output logits become less correlated as training progresses. This suggests that while the footstep candidate cost is useful for initial learning, the model may be locked into a local minimum where it relies too heavily on this feature.




Figure 4.11: Correlation between negative footstep candidate cost and GaitNet output logits during training. Negative costs are used so that a costmap which perfectly captures optimal footstep locations has a correlation of +1.

Now, we train the Cost-Ablated-GaitNet model, over-writing all footstep candidate cost values with zero. The training process is identical to that of the standard GaitNet model, ensuring a fair comparison. The performance of both models is then evaluated using the same metrics outlined in ??.




Figure 4.12: Evaluation of Cost-Ablated-GaitNet across various terrain difficulties and commanded velocities. Overall survival rate of 77.7%. Mean success rate measured as the percentage of 50 episodes which completed 20s without terminating, under the termination conditions described in ??.

?? presents the performance of the Cost-Ablated-GaitNet. The results indicates a measurable performance over the standard GaitNet formulation. GaitNet was able to achieve an overall survival rate of 69.4% (??), while the Cost-Ablated-GaitNet achieved a survival rate of 77.7% (??).

Looking at the mean episode reward in ??, we see a similar trend. The Cost-Ablated-GaitNet achieves a higher mean episode reward during training compared to the standard GaitNet. This does come at the cost of initially slower learning, where the Cost-Ablated-GaitNet trails the performance of the standard GaitNet for the first ~ 75 iterations.




Figure 4.13: Mean episode reward during training for GaitNet and Cost-Ablated-GaitNet.

Interestingly, the Cost-Ablated-GaitNet also demonstrates a slower cadence during training, as shown in ??.




Figure 4.14: Comparison of steps per second during training for GaitNet and Cost-Ablated-GaitNet.

Looking into how this reduced cadence affects the swing schedule, we see in ?? that under the same test parameters (10% terrain difficulty and 0.15 m/s commanded velocity), the Cost-Ablated-GaitNet takes less redundant steps—when the same foot is moved multiple times in succession.




Figure 4.15: Comparison of swing schedules for GaitNet (top) and Cost-Ablated-GaitNet (bottom) at 10% terrain difficulty with a commanded velocity of 0.15 m/s.

4.4.1 Discussion

The action cost ablation study provides perhaps the most significant insight. Counterintuitively, removing the footstep candidate cost from GaitNet’s input improves performance from 69.4% to 77.7% survival rate. This result challenges the initial assumption that providing the learned heuristic costs would accelerate training and improve final performance. Several factors likely contribute to this outcome. The correlation analysis (??) shows that while GaitNet initially relies heavily on the cost term (correlation near 0.8), this dependence decreases during training, suggesting the network learns to prioritize other state features.

The inclusion of the cost term may constrain exploration, biasing the network toward locally optimal solutions that align with the footstep evaluation network’s heuristics but prevent discovery of globally better strategies.

Intuitively this makes sense, as the footstep cost term included in the standard GaitNet model provides information about the robot dynamics, learned from the Footstep Evaluation Network. Without this information, the Cost-Ablated-GaitNet must learn the robot dynamics from scratch, leading to slower initial learning. However, as noted, this allows for more freedom to explore. Moreover, Cost-Ablated-GaitNet exhibits a naturally slower cadence (??) and fewer redundant footsteps (??), indicating more deliberate and efficient gait patterns. This behavior emerges without different reward shaping, suggesting that the model better captures the underlying dynamics when not constrained by pre-computed costs. The slower initial learning observed in Cost-Ablated-GaitNet represents a worthwhile trade-off for superior final performance.

4.5 General Synthesis of Results

These findings collectively suggest that GaitNet’s strength lies in its ability to learn coordinated, multi-leg motion strategies through reinforcement learning. The hierarchical architecture successfully constrains the action space to valid movements while allowing sufficient flexibility for dynamic gait generation. However, the system appears to benefit more from direct exploration of state-action relationships than from intermediate learned heuristics. This insight has important implications for hybrid control architectures: while learned perception modules can effectively filter candidate actions based on geometric and kinematic constraints, value estimation may be better left to end-to-end learning that captures the full dynamics of the system.

The results also highlight opportunities for improvement. The relatively modest absolute performance (77.7% best-case survival rate) indicates room for enhancement through reward function refinement, network architecture modifications, or improved low level control. Future work incorporating prediction horizons or more complex terrain may better exploit these capabilities. Overall, this work demonstrates that greedy, neural network-based planners can generate effective dynamic gaits for quadruped robots.

5 Conclusions and Future Work

5.1 Summary of Findings

This thesis presents GaitNet, a novel neural network-based greedy planner for generating dynamic, acyclic gaits in quadruped robots operating in challenging environments. Through a hierarchical hybrid control architecture that combines learned footstep evaluation with reinforcement learning-based gait selection, this work demonstrates that efficient, terrain-adaptive locomotion can be achieved without relying on predefined gait patterns.

GaitNet successfully generates non-periodic, dynamically stable gaits that adapt to varying terrain conditions and commanded velocities. Unlike traditional approaches that rely on fixed gait patterns (trot, walk, etc.), the system learns to coordinate multi-leg motions with variable timing, moving up to two legs simultaneously when conditions permit. The generated swing schedules (??) demonstrate true acyclic behavior, with the robot adjusting its gait pattern in response to both terrain difficulty and commanded velocity without following repetitive cycles.

Comparative evaluation against a physical Legged Perceptive Stack baseline reveals substantial performance improvements. GaitNet achieves a higher survival rate across diverse unstructured hardware terrain difficulties. This improvement is particularly pronounced at higher speeds and on more challenging terrain, demonstrating the value of continuous output space optimization and dynamic gait coordination.

The projected gradient ascent formulation over the output space successfully generates high-quality footstep candidates without relying on discrete sampling grids. The system demonstrates strong generalization across diverse robot states, accurately converging onto optimal foothold regions even in challenging configurations.

GaitNet learns to modulate swing durations based on environmental conditions, with mean swing duration of 0.24 s and standard deviation of 0.034 s during evaluation. The system generally selects swing durations between 0.225 s and 0.25 s, but occasionally selects shorter swing durations when needed. However, the swing duration ablation study reveals that this feature provides negligible performance benefits in the tested environments, suggesting either that the current scenarios do not fully exploit this capability or that the reward structure does not adequately incentivize optimal duration selection.

Perhaps the most surprising finding emerges from the action cost ablation study, which demonstrates that removing the footstep candidate cost from GaitNet’s input actually improves performance. Cost-Ablated-GaitNet achieves a 77.7% survival rate, representing an 8.3 percentage point improvement over the standard formulation. The correlation analysis (??) shows that while GaitNet initially relies heavily on the provided costs, this dependence

decreases during training as the network learns to prioritize other state features.

Cost-Ablated-GaitNet exhibits naturally more efficient behavior, taking fewer redundant steps and maintaining a lower overall cadence (??) without explicit reward shaping for these characteristics. This emergent efficiency suggests that when freed from pre-computed heuristic constraints, the reinforcement learning process discovers more globally optimal locomotion strategies. The comparison of swing schedules (??) illustrates this efficiency, with fewer instances of the same foot being repositioned multiple times in succession.

The greedy planning approach enables real-time operation suitable for online control. Unlike MCTS-based or full trajectory optimization methods that require substantial computation for each decision, GaitNet performs a single forward pass through a relatively compact neural network to evaluate and rank discrete action candidates. This computational efficiency is achieved while maintaining robust performance.

The hierarchical design successfully bridges learning-based and model-based control paradigms. By decomposing the problem into footstep candidate generation and gait selection, the architecture provides interpretability and safety through explicit action filtering while maintaining the adaptability of learned policies. The strict constraints on possible actions prevent invalid or unsafe motions while allowing sufficient flexibility for dynamic, terrain-adaptive behavior.

In summary, this work validates the hypothesis that a greedy, neural network-based planner integrated within a hybrid control framework can generate dynamic, acyclic gaits for robust quadruped locomotion. The findings provide valuable insights into the design of hybrid locomotion controllers, particularly regarding the role of learned heuristics and the balance between constraint and flexibility in action selection. While absolute performance leaves room for improvement, the demonstrated capabilities and unexpected insights from ablation studies establish a foundation for future development of adaptive legged locomotion systems.

5.2 Limitations

While the results presented in this work are promising, several limitations should be acknowledged. First, the proposed framework has only been validated in simulation. Real-world deployment may introduce additional challenges, such as unmodeled dynamics, sensor noise, and hardware constraints, which could affect performance. Second, GaitNet operates in a greedy, no-lookahead fashion, limiting its effectiveness at higher speeds or in highly dynamic environments where predictive planning could improve stability and efficiency. Additionally, the system’s performance is constrained by the accuracy of the low-level controller in executing footstep placements. Errors in tracking or positioning can degrade the overall effectiveness of the generated gaits. Finally, GaitNet’s performance is inherently tied to the quality of footstep candidates produced by the footstep evaluation network; suboptimal candidate sampling may restrict both the diversity and effectiveness of the generated actions.

These limitations highlight clear directions for future research, including real-world testing on physical hardware, the incorporation of predictive planning strategies, and improvements to both candidate generation and low-level control precision.

5.3 Future Work

Several avenues exist to extend and improve the current continuous optimization framework.

Swing re-planning is another potential improvement. Currently, GaitNet does not re-plan during the swing phase. Incorporating real-time re-planning would allow the robot to adjust to unexpected disturbances or terrain changes, this would require a more robust low-level controller implementation capable of accurately tracking the modified footstep trajectories.

Long-horizon planning also presents an opportunity for enhancement. GaitNet is presently trained to generate actions for a single point in time, which limits the MPC’s ability to predict future contact states accurately; the MPC expects the robot to revert to a nominal stance after each swing. While this is acceptable at low speeds, it could become problematic at higher velocities, where predictive planning would improve stability and performance.

Currently, the MPC runs on the CPU, limiting the speed of reinforcement learning. Leveraging GPU-accelerated MPCs, as explored in recent works [?, ?], could significantly accelerate training.

GaitNet also selects only one action at a time, leading to slightly staggered swing start times. Initial attempts to select multiple simultaneous actions encountered difficulties with gradient flow and learning stability. Developing a reliable method for multi-action selection could enable more fluid and dynamic gait patterns, though it would require careful network and training design.

Finally, improvements to the action candidate sampler could enhance GaitNet’s performance. At present, action selection relies on sampling-based methods, which may not always identify the highest-quality actions. This is seen in ?? by the fact that GaitNet relies less and less on the footstep candidate cost as training progresses. Directly searching the action space using techniques such as projected gradient descent or other optimization strategies could improve solution quality at the expense of additional computation. Another possibility is to train a candidate selection network jointly with GaitNet to improve the diversity and relevance of sampled actions.

5.4 Final Remarks

This thesis set out to explore whether a hybrid control pipeline, combining a greedy neural network-based planner with a model-based controller, could generate dynamic and acyclic gaits for quadruped locomotion. Through the design, training, and evaluation of GaitNet, this work demonstrated that such an approach is not only feasible but can yield robust and adaptive locomotion behaviors in challenging simulated environments. The results validate the central hypothesis that hybrid architectures can bridge the gap between the adaptability of learning-based methods and the reliability of model-based control.

Beyond performance metrics, the findings carry broader implications for the design of locomotion systems. The success of the greedy planning strategy highlights the potential of lightweight, inference-efficient neural networks for real-time decision-making in high-dimensional control problems.

Equally important are the insights into learning dynamics and representation. The discovery that removing pre-computed cost terms improves performance suggests that hybrid

controllers should not merely layer learned modules on top of existing heuristics, but rather allow networks to develop their own internal representations of value and risk through exploration. As such, the future of legged locomotion control may depend less on handcrafted structure and more on architectures that enable learned coordination within well-defined safety boundaries.

In closing, this work contributes to the growing evidence that data-driven methods can coexist effectively with analytical control frameworks, each compensating for the other’s weaknesses. GaitNet provides a concrete example of how this integration can produce dynamic, adaptive, and interpretable motion strategies—paving the way for future quadruped systems capable of operating reliably in unstructured, real-world environments. Continued efforts toward real-hardware validation, richer environments, and predictive planning will further advance the vision of truly agile, autonomous legged locomotion.

A GaitNet Training Configuration

A.1 Reward Function Analysis

Designing the reward functions for training GaitNet (see ??) proved to be a challenging task. Various combinations of rewards often resulted in undesirable behaviors, such as the robot frequently lifting its feet while idle or dragging feet during locomotion. Achieving a natural gait required careful balancing, particularly between *op_reward* and *mdp.foot_slip_penalty*.

Function ¹	Weight	Description
mdp.is_alive	0.4	Reward for being alive
mdp.track_lin_vel_xy_exp	0.5	Reward for tracking linear velocity in the XY plane
mdp.track_ang_vel_z_exp	0.5	Reward for tracking angular velocity around the Z axis
op_reward	-0.1	Penalty for performing actions
mdp.is_terminated	-200	Penalty for termination
mdp.lin_vel_z_l2	-2.5	Penalty for linear velocity in the Z direction
mdp.ang_vel_xy_l2	-0.1	Penalty for angular velocity in the XY plane
mdp.flat_orientation_l2	-8	Penalty for non-flat orientation
mdp.foot_slip_penalty	-6	Penalty for foot slip

Table A.1: Reward functions used to train GaitNet.

In addition to the rewards listed in Table ??, several other reward functions were explored but ultimately not used.

The *a_foot_in_swing* reward was intended to encourage the agent to frequently lift its feet during early learning. However, it proved unnecessary, as the initial network weights already favored foot movement. Furthermore, including this reward caused the agent to move its feet excessively when idle.

The *no_op_reward* was designed to encourage the agent to remain idle when no useful actions were available. In practice, this reward required an excessively high weighting to have any effect, at which point it would overshadow other rewards. The *op_reward* was found to be a more effective alternative for achieving the desired behavior.

¹Functions named "mdp.*" are built-in functions provided by the NVIDIA Isaac Lab framework.

A.2 Termination Functions

The termination functions used in training GaitNet are summarized in Table ?? . These functions help define the conditions under which an episode ends, either applying the termination penalty or simply signaling an invalid state.

Function ²	Time Out ³	Description
mdp.time_out	True	Terminate at the end of the episode
mdp.bad_orientation	False	Terminate if the robot’s orientation is too far from upright
mdp.root_height_below_minimum	False	Terminate if the robot’s base height is too low
mdp.terrain_out_of_bounds	True	Terminate if the robot leaves the terrain bounds
foot_in_void	False	Terminate if any foot steps into the void

Table A.2: Termination functions used to train GaitNet.

A.3 Commands

The command used in training GaitNet is summarized below. This command defines the highest level input to the environment, **u**.

- **UniformVelocityCommand**: This command samples a desired velocity vector $\mathbf{v} = [v_x, v_y, \omega_z]$ from a uniform distribution.
 - resampling time range: (2.5, 10) s
 - v_x range: (-0.2, 0.2) m/s
 - v_y range: (-0.2, 0.2) m/s
 - ω_z range: (-0.4, 0.4) rad/s
 - probability of zero command: 0.05

A.4 PPO Hyperparameters

The PPO hyperparameters used for training GaitNet are summarized in Table ?? . While not all parameters were rigorously tuned, they were found to perform well in practice. The discount factor *gamma* was specifically selected in relation to the agent observation frequency of 25,Hz to provide a reasonable effective horizon. The *learning_rate* was chosen relatively high to compensate for the slower data collection speed.

²Functions named "mdp.*" are built-in functions provided by the NVIDIA Isaac Lab framework.

³Time Out indicates whether the termination applies the mdp.is.terminated penalty.

Hyperparameter	Value
clip_param	0.3
num_learning_epochs	8
num_mini_batches	4
value_loss_coef	0.5
entropy_coef	0.02
learning_rate	3e-4
max_grad_norm	1.0
use_clipped_value_loss	True
gamma	0.99
lam	0.95

Table A.3: Hyperparameters used for PPO training.